



CONTENTS

1	Data tables in SQL	3
1.1	Using SQL from Python	6
2	Diodes, Rectifiers, and Photovoltaics	9
2.1	Silicon	9
2.2	Diodes	11
2.2.1	Rectifiers	14
2.2.2	Photovoltaics	14
3	Trees	17
3.1	Binary Trees and Binary Search Trees	18
3.2	File Structures and File System Hierarchies	19
3.3	Red-Black Trees and AVL Trees	21
3.4	B-Trees, Tries, and Suffix Trees	21
4	Tree Traversal	23
4.1	BFS	23
4.2	DFS	26
4.2.1	Pre-order	26
4.2.2	Post-order	27
4.2.3	In-order	27
4.3	Pseudocode	28
4.4	Recursive vs Iterative Traversal	28
4.4.1	Pre-order and Post-order Numbering	28

5	Finding Shortest Paths	29
6	Dijkstra's Algorithm	31
6.1	Algorithm Description	31
6.2	Implementation	33
6.3	Making it faster	37
7	Priority Queues	39
8	Binary Search	41
8.1	A Naive Implementation of the Priority Queue	41
8.2	Using the Priority Queue	42
8.3	Binary Search	44
8.4	Algorithm	45
9	Other Graph Algorithms and Proofs	47
9.1	Bellman-Ford Algorithm	47
9.2	Prim's Algorithm	47
9.3	Breadth-First Search	48
A	Answers to Exercises	49
	Index	51

Data tables in SQL

Most organizations keep their data as tables inside a relational database management system (compared to pandas, CSV, or spreadsheets). Developers talk to those systems using a language called SQL (“Structured Query Language”).

Some relational database managers are pricey products you may have heard of before, such as Oracle or Microsoft SQL Server. Some are free, such as PostgreSQL or MySQL. These are server software that client programs talk to over a company’s network.¹

There is a library, called `sqlite`, that lets us create files that hold tables. We can use SQL to create, edit, and browse those tables. `sqlite` is free, fast, and very easy to install. We will use `sqlite` instead of a networked database management system.

If you look in your digital resources, you will find a file called `bikes.db`. We created this file using `sqlite`, and now you will use `sqlite` to access it.

In the terminal, get to the directory where `bikes.db` lives. To open the `sqlite` tool on that file:

```
1 > sqlite3 bikes.db
```

(If your system complains that there is no `sqlite3` tool, you need to install `sqlite`. See this website: <https://sqlite.org/>)

Please follow along: type each command shown here into the terminal and see what happens.

We mostly run SQL commands in this tool, but there are a few non-SQL commands that all start with a period. To see the tables and their columns, you can run `.schema`:

```
1 sqlite> .schema
2 CREATE TABLE bike (bike_id int PRIMARY KEY, brand text, size int,
3     purchase_price real, purchase_date date, status text);
```

¹It should be noted that many notetaking and file storage applications allow you to run SQL-like queries to search your files to meet certain criteria. For example, Obsidian, the Markdown notetaking app, has a plugin called Dataview, which allows you to run searches for notes matching certain *metadata* attribute criteria. See more about this here: <https://github.com/blacksmithgu/obsidian-dataview>

That is the SQL command that we used to create the bike table. You can see all the columns and their types.

You want to see all the rows of data in that table?

```
1  sqlite> select * from bike;
2  4997391|GT|57|269.61|2009-05-03|rented
3  5429447|Cannondale|50|215.91|2002-02-17|broken
4  5019171|Trek|58|251.17|1985-07-11|rented
5  3000288|Cannondale|57|211.08|1993-01-05|broken
6  880965|GT|52|281.75|1995-08-02|available
7  ...
```

You will see 1000 rows of data!

The SQL language is not case-sensitive, so you can also write it like this:

```
1  sqlite> SELECT * FROM BIKE;
```

Often, you will see SQL with just the SQL keywords in all caps:

```
1  sqlite> SELECT * FROM bike;
```

The semicolon is not part of SQL, but it tells sqlite that you are done writing a command and that it should be executed.

SQL lets you choose which columns you would like to see. The asterisk (FIXME) used above signifies all columns, and the `bike_id, brand` only gets the bike's id and brand from the dataframe:

```
1  sqlite> SELECT bike_id, brand FROM bike;
2  4997391|GT
3  5429447|Cannondale
4  5019171|Trek
5  3000288|Cannondale
6  ...
```

Using WHERE, SQL lets you use conditions to decide which rows you would like to see, and can be combined with the common operators AND, OR, and NOT:

```
1  sqlite> SELECT * FROM bike WHERE purchase_date > '2009-01-01' AND brand = 'GT';
2  4997391|GT|57|269.61|2009-05-03|rented
3  326774|GT|56|165.0|2009-06-27|available
```

```

4 | 264933|GT|52|302.43|2009-07-09|available
5 | 5931243|GT|55|173.56|2009-11-26|rented
6 | 4819848|GT|51|221.71|2009-12-11|rented
7 | 9347713|GT|52|232.32|2009-06-13|available
8 | 3019205|GT|58|262.94|2009-08-22|available

```

Using DISTINCT, SQL lets you get just one copy of each value:

```

1 | sqlite> SELECT DISTINCT status FROM bike;
2 | rented
3 | broken
4 | available
5 |
6 | Busted
7 | Flat tire
8 | good
9 | out
10 | Rented

```

You can also edit these rows. For example, if you wanted every status that is Busted to be changed to broken, you can use an UPDATE statement with a SET:

```

1 | sqlite> UPDATE bike SET status='broken' WHERE status='Busted';
2 | sqlite> SELECT DISTINCT status FROM bike;
3 | rented
4 | broken
5 | available
6 | Flat tire
7 | good
8 | out
9 | Rented

```

You can insert new rows:

```

1 | sqlite> INSERT INTO bike (bike_id, brand, size, purchase_price, purchase_date,
2 |   status)
3 |   ...> VALUES (1, 'GT', 53, 123.45, '2020-11-13', 'available');
4 | sqlite> SELECT * FROM bike WHERE bike_id = 1;
5 | 1|GT|53|123.45|2020-11-13|available

```

Note that the bike_id here must be *unique*.

You can delete rows:

```
1 sqlite> DELETE FROM bike WHERE bike_id = 1;
2 sqlite> SELECT * FROM bike WHERE bike_id = 1;
```

To get out of sqlite, type `.exit`.

Exercise 1 SQL Query

Execute an SQL query that returns the `bike_id` (no other columns) of every Trek bike that cost more than \$300.

Working Space

Answer on Page 49

1.1 Using SQL from Python

The people behind sqlite created a library for Python that lets you execute SQL and fetch the results from inside a python program.

Let's create a simple program that fetches and displays the bike ID and purchase date of every Trek bike that cost more than \$300.

Create a file called `report.py`:

```
1 import sqlite3 as db
2
3 con = db.connect('bikes.db')
4 cur = con.cursor()
5
6 cur.execute("SELECT bike_id, purchase_date FROM bike WHERE purchase_price > 330
7             ↪ AND brand='Trek'")
8 rows = cur.fetchall()
9
10 today = datetime.date.today()
11 for row in rows:
12     print(f"Bike {row[0]}, purchased {row[1]}")
13
14 con.close()
```

When you execute it, you should see:

```
1 > python3 report.py
2 Bike 4128046, purchased 2007-08-06
3 Bike 7117808, purchased 1995-03-12
4 Bike 7176903, purchased 1986-07-03
5 Bike 827899, purchased 2009-03-14
6 Bike 363983, purchased 1970-08-16
```


Diodes, Rectifiers, and Photovoltaics

Early in this course, we discussed some basic electronics components: resistors, capacitors, and inductors. In this chapter, we are going to introduce a semiconductor: *Diodes* ensure that current flows in only one direction.

You will learn a lot about semiconductors from studying diodes. We can then apply what you have learned to solar cells, transistors, and integrated circuits like CPUs.

We are going to focus on semiconductors made of silicon. You can make semiconductors out of other materials, but silicon has proved to be very versatile and cost-effective.

2.1 Silicon

A silicon atom has 14 protons and 14 electrons. The inner 10 electrons are arranged in two complete and stable shells. The outer four electrons are in an incomplete shell, which allows them to form bonds with other atoms. We call these *valence electrons* because they aren't tightly bound to their atoms. A complete outer shell would have eight electrons.

Atoms can share electrons to get complete shells. This is called *covalent bonding*. When atoms share electrons, they form a *covalent bond*. The electrons are shared between the atoms, and the atoms are held together by the shared electrons.

In a silicon crystal, each silicon atom shares its outer four electrons with four neighboring atoms. With its four electrons and an electron borrowed from each of its four neighbors, each silicon atom has eight electrons in its outer shell. This is a complete shell, and the atoms are stable. This is why glass is hard and stable. It also why electricity doesn't flow through glass – all the electrons are tightly bound to their atoms.

Phosphorus has five valence electrons. If you put a phosphorus atom into silicon crystal, it shares four of its electrons with its neighbors, but it has one electron left over after the complete shell of eight is built. As you add more phosphorus to a silicon crystal, it conducts more electricity. We call this *doping*.




Figure 2.1: Phosphorus atom in a silicon crystal

Because this kind of doping frees up some electrons, which are negatively charged, we call this *n-doping*.

Boron has only three valence electrons. If you put a boron atom into silicon crystal, it shares three of its electrons with its neighbors, which leaves a "hole" that would like to be filled by an electron. We call this *p-doping*.




Figure 2.2: Boron atom in a silicon crystal

Both n-doped and p-doped silicon are conductors of electricity. In n-doped silicon, we say the loose electrons are the *carriers* of the current. In p-doped silicon, we say the holes are the carriers.

Note that doped silicon is still mostly silicon. In normal doping, maybe 1 in 10,000 atoms are replaced with boron or phosphorus.

2.2 Diodes

A diode is a semiconductor device that allows current to flow in one direction only.




Figure 2.3: Current flows in one direction through the diode, not in the other direction.

Diodes can be made with silicon that has been n-doped and p-doped. A diode looks like this:



Figure 2.4: What Diodes Look Like

We say that current flows from the positive terminal to the negative terminal of a battery. However, remember that the electrons are actually moving from the negative terminal to the positive terminal. This is a common source of confusion in any discussion of diodes.

One end of the diode is called the *anode*. The anode is p-doped. The other end is called the *cathode*. The cathode is n-doped.




Figure 2.5: A p-n Diode

The region close to where the anode and cathode meet is called the *depletion region*. In this space, the spare electrons in the cathode (the n-doped material) migrate over to fill the holes in anode (the p-doped material). This creates an electric field: there are more electrons than protons on the anode side, and more protons than electrons on the cathode side. (How strong is this field? It is typically about 0.7 volts.)




Figure 2.6: The depletion region in a diode

What happens if we overpower that 0.7 volts? What if we put 2 volts in the opposite direction? This shrinks the depletion region and the diode becomes a conductor. For example, if the negative end of a battery is connected to the cathode of a diode and the positive end is connected to the anode, current will flow through the diode.




Figure 2.7: Electrons flow through the diode

If we reverse the battery, the depletion region expands, and the diode becomes an insulator.




Figure 2.8: Reverse the electric field? No flow.

Here is the symbol for a diode:




Figure 2.9: Diode Symbol

Current will flow in the direction of the arrow. It will not flow in the opposite direction.

2.2.1 Rectifiers

One of the really useful things we can do with a diode is to convert alternating current (AC) into direct current (DC). This is called rectification. Remember that AC pushes and pulls the electrons back and forth in the wire. Rectification ensures that the electrons only move in one direction.

A *half-wave rectifier* uses a single diode to ensure that the current flows in one direction through the load.



Figure 2.10: Half-wave Rectifier

Current flows in the direction of the arrow on the diode. If the voltage is reversed, the diode will block the current from moving.

A *full-wave rectifier* uses four diodes to swap the direction of the current through the load when necessary.




Figure 2.11: Full-wave Rectifier

2.2.2 Photovoltaics

Photovoltaics (or "Solar cells") are similar to diodes. They have a p-n junction with a 0.7 volt electrical field. The difference is that sunlight is allowed to shine on the cathode

side. When an atom is struck by a photon of light energy, the electron becomes excited. Sometimes it is excited enough to leap across that 0.7 volt barrier.

When it leaps across, there is an extra electron on the anode side of the barrier and an extra proton on the cathode side. The electron is pushed back to where it came from, but it can't go back over the barrier (unless more than 0.7 volts of charge have built up across the barrier). So, it goes the long way: through the load.




Figure 2.12: Solar Cell

Trees

Trees are one of the most versatile and widely used data structures in computer science. A *tree* is a hierarchical data structure consisting of nodes, where each node has a value and a set of references (*edges*) to its child nodes. The node at the top of the hierarchy is called the *root*, and nodes with the same parent are called *siblings*, as seen in Figure 3.1.




Figure 3.1: A tree with a root node and 3 children.

to maintain its connected and acyclic properties. If there were more than $n - 1$ edges, it would create a cycle, violating the acyclic property. If there were fewer than $n - 1$ edges, it would not be fully connected, violating the connected property.

Less formally, a *forest* is a collection of multiple trees (may be disconnected), and a *tree* is a forest with only one tree. Lots of agricultural terms in this chapter!

Trees have a unique structural property: between any two nodes, there exists exactly one simple path. This follows directly from the fact that trees are both connected and acyclic. As a result, there is only one way to move from a start node to an end node.

Although trees are often drawn with a root at the top, this is a matter of convention. Any node in a tree can be chosen as the root, and doing so induces a hierarchical structure in which edges are interpreted as parent-child relationships.

Once a root is selected, every node (except the root) has exactly one parent, and zero or more children. Because of this relationship, any randomly selected node can form a tree (we call this a *subtree*, consisting of the node and all its descendants), making the tree structure *recursive*. Moving from one node to the next is the recursive step that minimizes

Trees are a subset of Graphs, a data structure comprised of nodes and edges. Graphs can be converted into trees by adding the following two constraints:

- The graph is connected, such that there exists a path between every pair of nodes
- The graph is acyclic; no cycles exist in the graph

With this comes the caveat that the number of edges in a tree is $\text{edges} = n - 1$, where n is the number of nodes in the tree. This is because a tree with n nodes must have exactly $n - 1$ edges

the size.

Nodes with no children are called, as you'd expect, *leaves*, as they are the outermost node of a tree. The *depth* of a chosen node is the number of edges from the *root* to that node, and the *height* of a tree is the maximum depth among all nodes. Equivalent to the height of a tree is the maximum number of edges from the *root* to any leaf. You may see the *level* of a node referred to as the number of edges from the *root* to that node, similar to depth. By convention, the level and depth of the root is 0. See Figure 3.2.

Some more relevant terminology you may encounter in the context of trees:

- *Path*: A sequence of nodes where each adjacent pair is connected by an edge. This can be used to show the idea of how to get between two nodes in a tree, for example, between two cities in a flight path network.
- *Ancestor*: A node that lies on the path from the root to a given node. This is done by repeatedly traversing from a child node to a parent node until the selected node is reached.
- *Descendant*: A node that lies on the path from a given node to a leaf. This is done by repeatedly traversing from a parent node to a child node until a leaf is reached.




Figure 3.2: A tree with a root node and 3 children, a few grandchildren, and 2 great grandchildren.

FIXME tree proof


3.1 Binary Trees and Binary Search Trees

The most common type of tree is the *binary tree*, where each node has *at most* two children, referred to as the *left child* and the *right child*. Binary trees are widely used in various applications, including expression parsing, binary search trees, and heaps.

Let's put this practice. Recall our binary search implementation from the recursion chapter. Given a chosen word and a sorted list of words, such as a dictionary, we can find the page of the chosen word by repeatedly dividing the available pages in half and comparing

the elements on the page to the chosen word. This is much easier than, for example, going through every single page and checking if the word is on that page.

Say we wanted to represent our dictionary as a tree that we could recursively binary search through. For simplicity, we will limit our words to mango, apple, pear, banana, anagram, peony and peach. These 7 words form a binary tree of height 2, as seen in Figure 3.3. We select the word mango as the root as it is the middle word after sorting alphabetically.



Similar to binary search, we can find any word in the tree by starting at the root and comparing the target word to the current node's word. If the target word is less than the current node's word, we move to the left child; if it is greater, we move to the right child. We repeat this process until we find the target word or reach a leaf node, which tells us that the target word is not in the tree. This structure allows us to cut the available words in half with each step in the tree we go.

FIXME expand on addition, removal, and search of binary trees
 FIXME add exercise of a similar dictionary tree with more words, have student add and remove words, and search for words

3.2 File Structures and File System Hierarchies

Computers, since their first creation in the 1980s, have needed a way to store and organize data. Prior to computers, many systems of organization were used to store physical data, such as filing cabinets, libraries, and warehouses.

These systems often had a hierarchical structure, with a top-level category that branched out into subcategories, with these subcategories further branching out into more specific categories, and so on.

For example, a library might have two top-level categories: fiction and non-fiction. Within the category non-fiction, you may find psychology, biographies, autobiographies, self-help, journalism books, or even nature guides. Under the category fiction, you may find romance, classics, crime, science fiction, thriller, or fantasy novels. Of course, these sub-genres can have even further categorizations too.

A hospital may have patient records categorized by first letter of last name, and then

sorted alphabetically from there.

With the development of the computer, sorting these types of hierarchical data became easy with the concept of directories. A *directory*, also called a *folder*, is a categorical way of organizing data (for example, by storing files or other subdirectories). We will stick to the more technical term of directories, however, any time you create a folder on your computer, you are making a *directory*. The name comes from how, in the twentieth-century, many phone books were referred to as telephone directories.


This structure naturally fits into the tree data structure:

- The **root** of the tree is the top-level directory (`/` on Linux/Unix¹ systems or `C:\` on Windows).
- A **node** is naturally fits the idea of either *directories nodes* or *files nodes*.
- A **directory node** may have children, where its children are subdirectories or files.
- A **file node** is a **leaf**, meaning it has no children.

The **path** to a file describes the unique sequence of directories from the root to that file. Test again test again test again

Example

A small file system might look like:



In tree terminology:

- `./` is the root node.

¹There is a very helpful video in Linux File Systems in your digital resources. Check it out!

- Downloads, Documents, and its subdirectories are internal nodes (directories).
- All files, such as those ending in .txt, .png, .exe, .csv, .docx, .py files are leaf nodes.

Side note: please do not store your passwords in anywhere on your computer's memory!

Exercise 2 File Paths

Working Space

In the above File Hierarchy Tree, what is the path to `data.csv`? What is the depth of the file tree at that point?

Answer on Page 49

3.3 Red-Black Trees and AVL Trees

3.4 B-Trees, Tries, and Suffix Trees

Tree Traversal

Trees can be “traversed” in a few different ways. When we say traversing a tree, we mean that we are visiting and extracting data from each node in a tree. We classify these *traversals* based on the order in which the nodes are visited.

This chapter will discuss *breadth-first search*, also known as *BFS*, which traverses in a *level-order*. Each level is traversed in sequential order.

Alternatively, there is *depth-first search* (*DFS*), which travels as far down a branch as possible before backtracking and checking other branches. Within DFS there are 3 common traversal methodologies: *in-order*, *pre-order* and *post-order*.

While the demonstrations for these traversals will be for trees, they work the same for any given graph, under one condition: using DFS or BFS on a *graph* requires you do have a preselected node a your *start node*.

4.1 BFS

Let’s look at running BFS on a tree. Starting out with this tree:




Figure 4.1: Our original tree containing a list of names to be traversed.

To perform BFS, we use a *queue*, which follows a first-in, first-out (*FIFO*) order. We begin by placing the root node into the queue. At each step, we remove the front node from the queue, visit it, and then add all of its children to the *back* of the queue.




Figure 4.2: Step 1: Visit Alan, then add its children (Brad, Connie) to the queue.

Here, we visit the root node first, which is Alan. After visiting Alan, we add its children, Brad and Connie, to the queue. Since there are no other nodes to visit at this level, we move on to the next level of the tree.




Figure 4.3: Step 2: Visit Brad, then add its children (Dean, Emily) to the queue.

We continue this process, visiting Connie next and adding its children, Faith and Gordon, to the queue. The queue now contains Connie, Dean, and Emily. Here is a step trace of the BFS traversal:

Step	Queue	Output (visited)
Start	[Alan]	[]
1	[Brad, Connie]	[Alan]
2	[Connie, Dean, Emily]	[Alan, Brad]
3	[Dean, Emily, Faith, Gordon]	[Alan, Brad, Connie]
4	[Emily, Faith, Gordon]	[Alan, Brad, Connie, Dean]

Continuing this process until the queue is empty gives the final traversal:

Alan, Brad, Connie, Dean, Emily, Faith, Gordon

This ordering reflects a level-by-level traversal of the tree, which is why BFS on a tree is also called level-order traversal.

To use BFS for *searching* rather than full traversal, we simply introduce a target value and stop the algorithm as soon as the target is found. Instead of visiting every node, we check each node when it is dequeued and terminate early if it matches the desired value.

For example, suppose we want to find “Faith” in the tree. BFS will examine nodes level by level: Alan, then Brad and Connie, then Dean, Emily, and Faith. As soon as Faith is encountered (we *dequeue* it from the queue), the search stops, without visiting any remaining nodes.

Our steps can be reflected as pseudocode as follows:

Algorithm 1 Breadth-First Search (BFS) - Iterative

```

procedure BFS(root, target)
  Q = empty queue
  ENQUEUE(Q, root)
  while Q is not empty do
    node = DEQUEUE(Q)
    if node = target then
      return node
    end if
    for each child of node do
      ENQUEUE(Q, child)
    end for
  end while
  return null
end procedure

```

This procedure is often used in game applications to find paths, such as maze solvers, chess algorithms, and tic-tac-toe solutions. It is also used in social network analysis to find the shortest path between two users, and in web crawling to explore the structure of

websites.

4.2 DFS

In contrast, we can run a *depth-first search* on the original tree (see Figure 4.1). Depth-first search, as the name implies, goes as deep down the tree until reaching a leaf, and then backtracks to the next possible, unsearched path. DFS (implicitly) uses a stack to explore the graph or tree.

The steps are fairly simple:

1. Push the root node into the stack.
2. Pop a node from the stack, and mark it as visited.
3. Push all unvisited adjacent nodes into the stack.
4. Repeat steps 2 and 3 until the stack is empty.

There are 3 orders in which we can traverse/search the deepest tree. We will look at pre-order, post-order, and in-order.

However, these orders essentially perform the same basic three tasks,

V Visit the selected node.

L Run DFS on the selected node's left subtree.

R Run DFS on the selected node's right subtree.

The recursion terminates when we reach a leaf. At this point, the algorithm returns to the previous call, which is how *backtracking* occurs. This backtracking is done via the stack we had mentioned previously.

4.2.1 Pre-order

In pre-order traversal, we visit the node *before* exploring its subtrees. The order is:

(V, L, R)

We visit the node *before* its children because the node's value is needed *prior* to exploring its subtrees. This is useful when the node represents a decision that affects which branch of the tree you go down, such as in tree copying or directory traversal like we talked about

in the previous chapter.

Using our original tree, the pre-order results are:

Alan, Brad, Dean, Emily, Connie, Faith, Gordon

4.2.2 Post-order

In post-order traversal, we visit the node *after* exploring both subtrees. The order is:

(L, R, V)

We visit the node *after* its children because the node's computation depends on results from both subtrees. This is useful when combining subtree results or deleting trees, where children must be handled before the parent.

The post-order traversal of our original tree is

Dean, Emily, Brad, Faith, Gordon, Connie, Alan

4.2.3 In-order

In in-order traversal, we visit the node *between* exploring its subtrees. The order is:

(L, V, R)

We visit the node *between* its subtrees because the correct interpretation of the node depends on partially processed children. In-order only makes sense for binary trees, because it returns the *sorted order* of the binary search tree. It depends on left and right subchildren. This would be like the alphabetical order of our dictionary example from the previous chapter.

The in-order traversal of our original tree becomes

Dean, Brad, Emily, Alan, Faith, Connie, Gordon

Algorithm 2 Depth-First Search (DFS) - Iterative

```
procedure DFS( $G, v$ )  
   $S$  = empty stack  
  PUSH( $S, v$ )  
  while  $S$  is not empty do  
     $v$  = POP( $S$ )  
    if  $v$  is not labeled as discovered then  
      label  $v$  as discovered  
      for each  $w$  adjacent to  $v$  do  
        PUSH( $S, w$ )  
      end for  
    end if  
  end while  
end procedure
```

4.3 Pseudocode

DFS is particularly useful for solving problems like connected-component detection in graphs and maze-solving.

4.4 Recursive vs Iterative Traversal

As we've established, trees are recursively defined structures, and as such, our algorithms for them can be defined as such too.

4.4.1 Pre-order and Post-order Numbering

Finding Shortest Paths

Dijkstra's Algorithm

Edsger W. Dijkstra was a great Dutch computer scientist. He came up with an algorithm for finding the cheapest path through a graph with weighted edges. Today it is known as *Dijkstra's Algorithm*. It is used in a wide variety of common problems, and is also both simple and elegant.

6.1 Algorithm Description

You are going to mark each node with how much it would cost to get there from some chosen origin node. For example, if you are shipping a container from Long Beach, you will mark each city with the cost of getting the container to that city.

You start by marking the price for Long Beach to zero (the container is already there). You then mark each adjacent city with the cost on the edge (figure 6.1a). Next, you declare Long Beach to be “visited” (figure 6.1b).




Figure 6.1: A map of weighted cities.

After that, you find the cheapest of the unvisited nodes. In this case, Los Angeles is

cheaper than San Diego, so that is the node you will visit next.

You mark all of the unvisited nodes adjacent to Los Angeles, with the price to ship it to Los Angeles plus the cost of shipping the container from Los Angeles to that city (figure 6.2a). Note that Bakersfield is marked with \$22.




Figure 6.2: Visiting the next two minimum cost cities.

Now, the cheapest unvisited node is San Diego, so you mark its neighbors with the cost to ship to San Diego, plus the price to ship from San Diego to the neighbor (figure 6.2b).

Notice that Bakersfield is already labeled with \$22 from a route through Los Angeles. But the price would be \$21 if you shipped it to Bakersfield via San Diego. Because the new route is cheaper, you change the price to the lower value.

(What does it mean that a node is “visited”? If a node is marked visited, it is marked with a price that won’t get any smaller.)

You continue visiting the cheapest unvisited node until all the nodes have been visited. Then you know every node has been marked with its lowest price.

In a big graph, each node may be marked several times in this process — each time with a lower price from a cheaper router.

Of course, once you have the price, you will ask, “What is the route that gets me that price?” So, we will also mark each node with the neighbor from which it would receive the shipment — the previous node. This is easy to do as we execute the algorithm.




Figure 6.3: Marking each node with a previous pointer.

Now, to figure out the cheapest route from Long Beach to Bakersfield, we start at the destination and follow the prev pointer back through San Diego, then to Long Beach.

Exercise 3 Expanding on the Delivery Graph

Notice that the path between Los Angeles to Bakersfield is not used once we add in our previous labels. Why might this be?

Working Space

Answer on Page 49

6.2 Implementation

We do not actually want to sully our graph objects with the three additional pieces of information we need:

- The current minimal cost from the origin node. This is usually called the *dist*, from “distance”.
- The neighbor who gives us the current minimal cost. This is usually called *prev*, from “previous”.
- Whether the city is visited or now.

So, we will keep them in collections external to the graph.

For example, to keep track of the *dist*, we will have a *dist* dictionary. Each node will be a key, the current minimal cost will be the value. If the node hasn't received even a first cost, we will put in infinity as the cost.

(After the algorithm is run, if the cost of a node is still infinity, that means that it cannot be reached from the origin node.)

We will also have a *prev* dictionary. The final node will be the key, and its previous neighbor will be the value.

Finally, the graph has a list of all the nodes, so we can just keep a set of the unvisited nodes.

Add a method to the *Graph* class that implements Dijkstra's algorithm:

```

1  def cost_from_node(self, origin_node):
2      # Cost of cheapest path from origin node discovered so far
3      # Initially the origin is zero and all the other are infinity
4      dist = {k: math.inf for k in self.nodes}
5      dist[origin_node] = 0.0
6
7      # The previous city on that cheapest path
8      prev = {}
9
10     # All the nodes start as unvisited
11     unvisited = set(self.nodes)
12
13     # While there are still unvisited nodes
14     while unvisited:
15
16         # Find unvisited node with lowest cost
17         min_cost = math.inf
18         for u in unvisited:
19             if dist[u] < min_cost:
20                 current_node = u
21                 min_cost = dist[u]
22
23         # If none are less than inf, we are done
24         # This happens in graphs that are not connected

```

```

25         if min_cost == math.inf:
26             return (dist, prev)
27
28         # Remove the lowest cost node from the unvisited list
29         unvisited.remove(current_node)
30
31         # Update all the unvisited neighbors
32         for edge in current_node.edges:
33
34             # What node is at the other end of this edge?
35             v = edge.other_end(current_node)
36
37             # Visited nodes are already minimized, skip them
38             if v not in unvisited:
39                 continue
40
41             # Is this a shorter route?
42             alt = dist[current_node] + edge.cost
43             if alt < dist[v]:
44
45                 # Update the distance and prev dicts
46                 dist[v] = alt
47                 prev[v] = current_node
48
49         return (dist, prev)

```

Append some code to your `cities.py` that tests this method:

```

1  (cost_from_long_beach, prev) = network.cost_from_node(long_beach)
2  print(f"\nMinimum costs from Long Beach = {cost_from_long_beach}")
3  print(f"\nLast city before = {prev}")
4
5  nyc_cost = cost_from_long_beach[nyc]
6
7  if nyc_cost < math.inf:
8      print(f"\n*** Total cost from Long Beach to NYC: ${nyc_cost:.2f} ***")
9  else:
10     print("You can't get to NYC from Long Beach")

```

When you run it, you should get a list of how much it costs to ship a container to each city from Long Beach:

```

1  Minimum costs from Long Beach = {(node:Long Beach, edges:1): 0.0,
2  (node:Los Angeles, edges:3): 12.0, (node:Denver, edges:3): 35.0,
3  (node:Phoenix, edges:3): 31.0, (node:Louisville, edges:3): 49.0,
4  (node:Cleveland, edges:4): 44.0, (node:Boston, edges:2): 57.0,
5  (node:New York City, edges:3): 52.0}

```

You will also get a collection of node pairs. What are these? For each node, you get the node that you would pass through on the cheapest route from Long Beach:

```

1 Last city before = {(node:Los Angeles, edges:3):(node:Long Beach, edges:1),
2 (node:Denver, edges:3):(node:Los Angeles, edges:3),
3 (node:Phoenix, edges:3):(node:Los Angeles, edges:3),
4 (node:Louisville, edges:3):(node:Denver, edges:3),
5 (node:Cleveland, edges:4):(node:Denver, edges:3),
6 (node:New York City, edges:3):(node:Cleveland, edges:4),
7 (node:Boston, edges:2):(node:Cleveland, edges:4)}

```

Your users will not want to read this; give them the shortest path as a list. Add a function to `graph.py` that turns the `prev` table into a path of nodes that lead from the origin to the destination:

```

1 def shortest_path(prev, destination):
2
3     # Include the destination in the path
4     path = [destination]
5     current_node = destination
6
7     # Keep stepping backward in the path
8     while current_node in prev:
9
10        # What node should come before the current node?
11        previous_node = prev[current_node]
12
13        # Insert it at the start of the list
14        path.insert(0, previous_node)
15        current_node = previous_node
16
17     return path

```

Test that out:

```

1 if nyc_cost < math.inf:
2     print(f"*** Total cost from Long Beach to NYC: ${nyc_cost:.2f} ***")
3
4     path_to_nyc = graph.shortest_path(prev, nyc)
5     print(f"*** Cheapest path from Long Beach to NYC: {path_to_nyc} ***")
6 else:
7     print("You can't get to NYC from Long Beach")

```

This should look like this:

```
1 *** Cheapest path from Long Beach to NYC: [(node:Long Beach, edges:1),  
2 (node:Los Angeles, edges:3), (node:Denver, edges:3), (node:Cleveland, edges:4),  
3 (node:New York City, edges:3)] ***
```

6.3 Making it faster

On really big networks, doing a full Dijkstra's algorithm would take too long; luckily, there are many methods for getting similar results quickly. When you ask for directions from Google Maps, it does not do a full Dijkstra's Algorithm for every possible route — it would just take too long.

However, there is a way to speed up this implementation. Look at this snippet:

```
1         # Find unvisited node with lowest cost  
2         min_cost = math.inf  
3         for u in unvisited:  
4             if dist[u] < min_cost:  
5                 current_node = u  
6                 min_cost = dist[u]
```

We are scanning through the list of all unvisited nodes, one by one, looking for the one with the lowest cost. If we kept this list sorted by cost, then the next one to visit would always be the first one in the list. This is done with a *priority queue* – a list that keeps itself sorted by some priority number – in this case the cost. In python, the standard priority queue is `heapq`.

(So why didn't I implement this using `heapq`? For Dijkstra's Algorithm, the nodes' priority – the current cost – changes as we find cheaper routes. `heapq` doesn't handle the changing priority very gracefully.)

In the next chapter, you will make a priority queue class that will work in this case.

Priority Queues

Binary Search

As mentioned in the last chapter, you are going to make a priority queue for use with Dijkstra's Algorithm. Using it will look like this algorithm called a *priority queue*:

```
1 import kqueue
2
3 myqueue = kqueue.PriorityQueue()
4 myqueue.add(long_beach, 0) # Inserts first city and its cost
5 myqueue.add(san_diego, 14) # Puts San Diego after Long Beach
6 myqueue.add(los_angeles, 12) # Inserts LA between Long Beach and San Diego
7 current_city = myqueue.pop() # Returns first city (Long Beach) and removes it
```

Now, if an item gets a new priority, we need to remove it and reinsert it in the new spot.

```
1 myqueue.add(city_a, 16) # Puts it last in the queue
2 myqueue.update(city_a, 16, 13) # Moves it to between LA and San Diego
```

8.1 A Naive Implementation of the Priority Queue

Create a file called `kqueue.py`. Let's do a simple implementation that stores the priority and the data as tuple. And we will keep it sorted by the priority. If two tuples have the same priority, we will sort by the data.

Type this in to `kqueue.py`:

```
1 class PriorityQueue:
2     def __init__(self):
3         self.list = []
4
5     # Return and remove the first item
6     def pop(self):
7         if len(self.list) > 0:
8             return self.list.pop(0)
9         else:
10            return None
11
12     def __len__(self):
13         return len(self.list)
```

```

14
15     def update(self, value, old_priority, new_priority):
16         old_pair = (old_priority, value)
17         self.list.remove(old_pair)
18         self.add(value, new_priority)
19
20     def add(self, value, priority):
21         pair = (priority, value)
22         # Add it at the end
23         self.list.append(pair)
24         # Resort the list
25         self.list.sort()

```

This will work fine, but it could be much more efficient:

- Every time we add a single element, we resort the whole list.
- The function `remove` is searching the list sequentially for the item to delete.

In a minute, we will revisit these inefficiencies and make them better.

8.2 Using the Priority Queue

We are going to change `graph.py` to use the priority queue. While we are doing so, why don't we also shrink the memory footprint of our program a bit.

Notice that as the algorithm is running, each node is in one of three states:

- Unseen: In the earlier implementation, these were the nodes with `math.inf` as their cost.
- Seen, but not finalized: These are “unvisited”, but don't have `math.inf` as their cost.
- Finalized: These are the “visited” nodes — we know that their cost won't decrease any more.

We can shrink the memory foot print by not putting the unseen into the `dist` dictionary at all. And instead of a separate set for “unvisited”, what if we moved finalized nodes and their distances into a separate dictionary?

Rewrite the `cost_from_node` function in `graph.py`:

```

1         # Visited nodes are already minimized, skip them
2         if v in finalized_dist:
3             continue

```

```

4
5         # What is the cost to this neighbor?
6         alt = current_node_cost + edge.cost
7
8         # Is this the first time I am seeing the node?
9         if v not in seen_dist:
10
11             # Insert into the seen_dict, prev, and priority queue
12             seen_dist[v] = alt
13             prev[v] = current_node
14             pqueue.add(v, alt)
15
16         else: # v has been seen. Is this a cheaper route?
17             old_dist = seen_dist[v]
18             if alt < old_dist:
19                 # Update the seen_dict, prev, and priority queue
20                 seen_dist[v] = alt
21                 prev[v] = current_node
22                 pqueue.update(v, old_dist, alt)
23
24         return (finalized_dist, prev)

```

This should have exactly the same result, except for the unreachable nodes. If you have a graph that is not connected, there will be nodes that cannot be reached from the origin. In the old version, these had a cost of `math.inf`. Now, they will just not be in the dictionary at all. So, change `cities.py` to deal with this:

```

1 if nyc in cost_from_long_beach:
2     nyc_cost = cost_from_long_beach[nyc]
3     print(f"\n*** Total cost from Long Beach to NYC: ${nyc_cost:.2f} ***")
4
5     path_to_nyc = graph.shortest_path(prev, nyc)
6     print(f"\n*** Cheapest path from Long Beach to NYC: {path_to_nyc} ***")
7 else:
8     print("You can't get to NYC from Long Beach")

```

If you run `cities.py` now, it should behave exactly like the old version.

However, there is a bug. It will rear its head if two cities with the same cost are in the priority queue together. Change `cities.py` so that Denver and Phoenix have the same cost:

```

1 graph.WeightedEdge(12, long_beach, los_angeles)
2 graph.WeightedEdge(19, los_angeles, denver)
3 graph.WeightedEdge(19, los_angeles, phoenix)

```

Now try running it. You should get an error:

```
1 TypeError: '<' not supported between instances of 'Node' and 'Node'
```

What happened? The `loc_for_pair` method is comparing tuples made up of a float and a `Node`. The float comes first in the tuple, so that is compared first. However, if the two tuples have the same priority, it then compares nodes.

The error statement says, “Nodes don’t have a less-than method; I don’t know how to compare them.”

Each `Node` lives at an address in memory. You can get that address as a number using the `id` function. The ID is unique and constant over the life of the object. It is a rather arbitrary ordering, but it will work for this problem. Add a method to your `Node` class:

```
1 # Nodes will be ordered by their location in memory
2 def __lt__(self, other):
3     return id(self) < id(other)
```

Fixed.

Now, let’s make the priority queue more efficient.

8.3 Binary Search

The phone company in every town used to print a thing called a phone book. The names and phone numbers were arranged alphabetically. As you might imagine, these books often had more than a thousand pages.

If you were looking for “John Jeffers”, you would not start at the first page and read sequentially until you reached his name. You would open the book in the middle, and see a name like “Mac Miller”, then think “Jeffers comes before Miller”. You would then split the pages in your left hand in half and see a name like “Hester Hamburg” and think “Jeffers comes after Hamburg”. You would then split the pages in your right hand, and continue on like this until you found the page with “John Jeffers” on it. That is binary search.

Binary Search is a search algorithm that finds the position of a target value within a sorted array. The binary search algorithm works by repeatedly dividing the search interval in half. If the target value is equal to the middle element of the array, the position is returned. If the target value is less than or greater than the middle element, the search continues in the lower or upper half of the array, respectively.

8.4 Algorithm

The binary search algorithm can be described as follows:

1. If the array is empty, the search is unsuccessful, so return "Not Found".
2. Otherwise, compare the target value to the middle element of the array.
3. If the target value matches the middle element, return the middle index.
4. If the target value is less than the middle element, repeat the search with the lower half of the array.
5. If the target value is greater than the middle element, repeat the search with the upper half of the array.
6. Repeat steps 2-5 until the target value is found or the array is exhausted.

```

1  class PriorityQueue:
2      def __init__(self):
3          self.list = []
4
5      # Return and remove the first item
6      def pop(self):
7          if len(self.list) > 0:
8              return self.list.pop(0)
9          else:
10             return None
11
12     def __len__(self):
13         return len(self.list)
14
15     def add(self, value, priority):
16         pair = (priority, value)
17         i = self.loc_for_pair(pair)
18         self.list.insert(i, pair)
19
20     def update(self, value, old_priority, new_priority):
21         old_pair = (old_priority, value)
22         i = self.loc_for_pair(old_pair)
23         del self.list[i]
24         self.add(value, new_priority)
25
26     def loc_for_pair(self, pair):
27         # The range where it could be is [lower, upper)
28         # Start with the whole list
29         lower = 0
30         upper = len(self.list)
31
32         while upper > lower:
33             next_split = (upper + lower) // 2

```

```
34         v = self.list[next_split]
35         if pair < v: # pair is to the left
36             upper = next_split
37         elif pair > v: # pair is to the right
38             lower = next_split + 1
39         else: # Found pair!
40             return next_split
41     return lower
```

If you try running it now, it should work perfectly.

You now have a graph class that will find the cheapest path quickly, even if it has thousands of nodes with thousands of edges.

Other Graph Algorithms and Proofs

Now that you are familiar with Dijkstra's algorithm for finding the shortest path in a graph, you are well-equipped to understand more graph algorithms. This document will discuss two other important algorithms: *Depth-First Search* (DFS) and the Bellman-Ford algorithm.

9.1 Bellman-Ford Algorithm

The Bellman-Ford Algorithm is another shortest path algorithm like Dijkstra's. However, unlike Dijkstra's Algorithm, Bellman-Ford can handle graphs with negative weight edges.

The algorithm works as follows:

1. Assign a tentative distance value for every node: set it to zero for our initial node and to infinity for all other nodes.
2. For each edge (u, v) with weight w , if the current distance to v is greater than the distance to u plus w , update the distance to v to be the distance to u plus w .
3. Repeat the previous step $|V| - 1$ times, where $|V|$ is the number of vertices in the graph.
4. After the above steps, if you can still find a shorter path, there exists a negative cycle.

If the graph does not contain a negative cycle reachable from the source, the shortest paths are well-defined, and Bellman-Ford will correctly calculate them. If a negative cycle is reachable, no solution exists, but Bellman-Ford will detect it.

9.2 Prim's Algorithm

Prim's Algorithm is minimum spanning tree FIXME <https://github.com/KontinuaFoundation/sequence/issues/315>

9.3 Breadth-First Search

Answers to Exercises

Answer to Exercise 1 (on page 6)

```
1 SELECT bike_id FROM bike WHERE purchase_price > 330 AND brand='Trek'
```

Answer to Exercise 2 (on page 21)

The path to `data.csv` is

```
./Documents/Projects/project2/data.csv
```

The depth is the number of edges from the root to the node.

In this case, it is the number of subdirectories needed to reach `data.csv`:

- (1) `./` → `Documents`
- (2) `Documents` → `Projects`
- (3) `Projects` → `project2`
- (4) `project2` → `data.csv`

So `data.csv` is at a depth of 4.

Answer to Exercise 3 (on page 33)

Los Angeles → Bakersfield is not chosen because it does not give the best known cost to Bakersfield. The `prev` pointer for Bakersfield goes to San Diego instead, since $\$21 < \22 .

The tree that is generated is called a *predecessor tree* or *explored tree* of the resulting edges formed by only the previous edges. We can use this to track shortest paths from a given origin node.



INDEX

Ancestor, 18
anode, 11

Bellman-Ford algorithm, 47
BFS, 23
Binary Search, 44
binary tree, 18
boron, 10
breadth-first search, 23

carriers, 10
cathode, 11
covalent bond, 9
covalent bonding, 9

depletion region, 12
depth, 18
Depth-First Search, 47
depth-first search, 23, 47
Descendant, 18
DFS, 23, 47
Dijkstra's Algorithm, 31
Dijkstra, Edsger, 31
diode, 9
Diodes, 9
directory, 20
doping, 9

edges, 17
explored tree, 49
FIFO, 24

folder, 20
forest, 17
full-wave rectifier, 14

half-wave rectifier, 14
height, 18

in-order, 23

leaves, 18
left child, 18
level, 18
level-order, 23

n-doped, 10
n-doping, 10

p-doped, 10
p-doping, 10
Path, 18
phosphorus, 9
photovoltaics, 15
post-order, 23
pre-order, 23
predecessor tree, 49
priority queue, 37, 41

queue, 24

rectifier, 14
recursive, 17
right child, 18

root, [17](#)

siblings, [17](#)

silicon, [9](#)

solar cells, [15](#)

SQL, [3](#)

SQLite, [3](#)

subtree, [17](#)

traversals, [23](#)

tree, [17](#)

valence electrons, [9](#)